

# Image classifier for the SVHN dataset

```
In [1]: import tensorflow as tf
        from scipy.io import loadmat

        import numpy as np
        import matplotlib.pyplot as plt
        %matplotlib inline
```

```
In [2]: # Loading the dataset

        train = loadmat('train_32x32.mat')
        test = loadmat('test_32x32.mat')
```

```
In [3]: # Inspecting the dataset

        print(train.keys())
        print(train['X'].shape)
        print(np.unique(train['y']))
        print(np.unique(test['y']))

dict_keys(['__header__', '__version__', '__globals__', 'X', 'y'])
(32, 32, 3, 73257)
[ 1  2  3  4  5  6  7  8  9 10]
[ 1  2  3  4  5  6  7  8  9 10]
```

```
In [4]: #Divided images values by 255 so they fall between 0 and 1
        train_labels = train['y']
        train_images = train['X']/255
        test_labels = test['y']
        test_images = test['X']/255
```

```
In [5]: #Labels that are the number 10 correspond to training images of a zero. The
        #the output layer softmax in the model won't work

        train_labels[train_labels[:,]==10] = 0
        test_labels[test_labels[:,]==10] = 0
```

```
In [6]: #Shuffle images
        def shuffle(images, labels, m):
            random_index = list(np.random.permutation(m))
            random_im = images[:, :, :, random_index]

            random_labels = labels[random_index]

            return random_im, random_labels

        train_images, train_labels = shuffle(train_images, train_labels, train_images.shape[0])
```

In [7]:

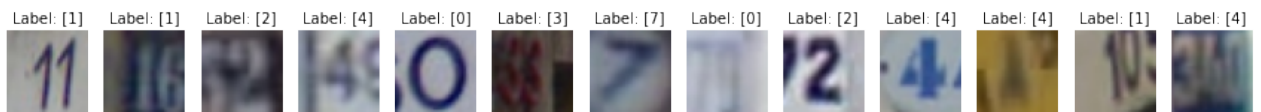
```
# Select random training images to display

train_examples_num = train_images.shape[3] #number of train examples
random_ind = np.random.choice(train_examples_num, 13) #generate 13 numbers
random_train_im = train_images[:, :, :, random_ind] #take 13 random images

random_train_lab = train_labels[random_ind] #take corresponding labels

fig, axes = plt.subplots(nrows=1, ncols=13, figsize=(15, 15)) #generate subplots

for (i, ax) in enumerate(axes.flat): # flatten the generated axes
    img = random_train_im[:, :, :, i] #take random image i
    ax.set_axis_off()
    ax.imshow(np.squeeze(img))
    ax.set_title('Label: ' + str(random_train_lab[i]), fontsize=10)
```



In [8]:

```
#Convert images to gray

def convert_gray(data, labels, print_it):
    num_toprint=13
    m = data.shape[3] #number of examples
    data = np.average(data, axis = 2)

    if print_it:
        ones = np.ones((32, 32, 3, 1))

        random_ind = np.random.choice(m, num_toprint) #random index to print
        random_data_gr = data[:, :, :, random_ind] #take images to print
        random_lab = labels[random_ind] #take labels

        print_data = random_data_gr[:, :, np.newaxis, :] * ones #add dummy inc
        print_labels = random_lab

        fig, axes = plt.subplots(nrows=1, ncols=num_toprint, figsize=(15, 15))
        for (i, ax) in enumerate(axes.flat):

            img = print_data[:, :, :, i]
            ax.set_axis_off()
            ax.imshow(np.squeeze(img))
            ax.set_title('Label: ' + str(print_labels[i]), fontsize=10)

    return data.T

train_images = convert_gray(train_images, train_labels, True)
test_images_gr = convert_gray(test_images, test_labels, False)
```



## MLP neural network classifier

```
In [9]: from tensorflow.keras.models import Sequential, load_model
from tensorflow.keras.layers import Dense, Flatten, Softmax, Conv2D, MaxPool2D
from tensorflow.keras.callbacks import ModelCheckpoint, Callback, CSVLogger
import pandas as pd
```

```
In [10]: def get_model(input_shape):

    model = Sequential([
        Flatten(input_shape=input_shape),
        Dense(128, activation='tanh', kernel_initializer='he_normal', kernel_regularizer=tf.keras.regularizers.l2(0.001)),
        Dense(64, activation='tanh', kernel_initializer='he_normal', bias_initializer='zeros', kernel_regularizer=tf.keras.regularizers.l2(0.001)),
        Dense(64, activation='tanh', kernel_initializer='he_normal', bias_initializer='zeros', kernel_regularizer=tf.keras.regularizers.l2(0.001)),

        Dense(10, activation='softmax')
    ])
    return model
```

```
In [11]: print(train_images[0,:,:].shape)

(32, 32)
```

```
In [12]: model = get_model(train_images[0,:,:].shape)
print(model.summary())
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 128)	131200
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 64)	4160
dense_3 (Dense)	(None, 10)	650
=====		
Total params: 144,266		
Trainable params: 144,266		
Non-trainable params: 0		

None

In [13]:

```
model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.0001),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy'])
print(train_images[0,:,:].shape)
print(train_labels.shape)
print(train_images.shape)

checkpoint = ModelCheckpoint('Model_1', save_weights_only=False, save_best_

history = model.fit(train_images,train_labels, epochs=30, batch_size=10, va
                    callbacks=[checkpoint, tf.keras.callbacks.CSVLogger("Mo

(32, 32)
(73257, 1)
(73257, 32, 32)
Epoch 1/30
6227/6227 [=====] - 10s 2ms/step - loss: 2.4780 - 
accuracy: 0.2381 - val_loss: 2.1732 - val_accuracy: 0.3587
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 2/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.8800 - a
ccuracy: 0.4658 - val_loss: 1.7399 - val_accuracy: 0.5159
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 3/30
6227/6227 [=====] - 9s 2ms/step - loss: 1.5959 - a
ccuracy: 0.5755 - val_loss: 1.5313 - val_accuracy: 0.6016
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 4/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.4627 - a
ccuracy: 0.6268 - val_loss: 1.4372 - val_accuracy: 0.6366
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 5/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.3783 - a
ccuracy: 0.6539 - val_loss: 1.3547 - val_accuracy: 0.6628
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 6/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.3139 - a
ccuracy: 0.6752 - val_loss: 1.3294 - val_accuracy: 0.6668
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 7/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.2690 - a
ccuracy: 0.6860 - val_loss: 1.2705 - val_accuracy: 0.6827
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 8/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.2293 - a
ccuracy: 0.6995 - val_loss: 1.2133 - val_accuracy: 0.7030
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 9/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.1962 - a
ccuracy: 0.7089 - val_loss: 1.2029 - val_accuracy: 0.7022
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 10/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.1687 - a
ccuracy: 0.7160 - val_loss: 1.1865 - val_accuracy: 0.7018
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 11/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.1466 - a
ccuracy: 0.7212 - val_loss: 1.1750 - val_accuracy: 0.7062
```

```
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 12/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.1270 - accuracy: 0.7272 - val_loss: 1.1579 - val_accuracy: 0.7108
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 13/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.1092 - accuracy: 0.7308 - val_loss: 1.1376 - val_accuracy: 0.7210
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 14/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.0919 - accuracy: 0.7363 - val_loss: 1.1155 - val_accuracy: 0.7294
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 15/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.0795 - accuracy: 0.7387 - val_loss: 1.0809 - val_accuracy: 0.7371
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 16/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.0638 - accuracy: 0.7421 - val_loss: 1.1114 - val_accuracy: 0.7265
Epoch 17/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.0498 - accuracy: 0.7489 - val_loss: 1.0616 - val_accuracy: 0.7431
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 18/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.0418 - accuracy: 0.7489 - val_loss: 1.0930 - val_accuracy: 0.7303
Epoch 19/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.0309 - accuracy: 0.7512 - val_loss: 1.1088 - val_accuracy: 0.7257
Epoch 20/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.0212 - accuracy: 0.7544 - val_loss: 1.0407 - val_accuracy: 0.7460
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 21/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.0131 - accuracy: 0.7571 - val_loss: 1.0769 - val_accuracy: 0.7372
Epoch 22/30
6227/6227 [=====] - 9s 1ms/step - loss: 1.0024 - accuracy: 0.7609 - val_loss: 1.0115 - val_accuracy: 0.7604
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 23/30
6227/6227 [=====] - 9s 1ms/step - loss: 0.9946 - accuracy: 0.7637 - val_loss: 1.0201 - val_accuracy: 0.7531
Epoch 24/30
6227/6227 [=====] - 9s 1ms/step - loss: 0.9871 - accuracy: 0.7658 - val_loss: 1.0356 - val_accuracy: 0.7478
Epoch 25/30
6227/6227 [=====] - 9s 2ms/step - loss: 0.9787 - accuracy: 0.7673 - val_loss: 1.0219 - val_accuracy: 0.7535
Epoch 26/30
6227/6227 [=====] - 10s 2ms/step - loss: 0.9729 - accuracy: 0.7697 - val_loss: 1.0031 - val_accuracy: 0.7580
INFO:tensorflow:Assets written to: Model_1/assets
Epoch 27/30
6227/6227 [=====] - 9s 1ms/step - loss: 0.9652 - accuracy: 0.7709 - val_loss: 1.0199 - val_accuracy: 0.7542
Epoch 28/30
6227/6227 [=====] - 9s 1ms/step - loss: 0.9572 - accuracy: 0.7746 - val_loss: 0.9915 - val_accuracy: 0.7643
INFO:tensorflow:Assets written to: Model_1/assets
```

```
Epoch 29/30
6227/6227 [=====] - 9s 1ms/step - loss: 0.9522 - a
ccuracy: 0.7754 - val_loss: 1.0023 - val_accuracy: 0.7576
Epoch 30/30
6227/6227 [=====] - 9s 1ms/step - loss: 0.9473 - a
ccuracy: 0.7783 - val_loss: 0.9778 - val_accuracy: 0.7640
INFO:tensorflow:Assets written to: Model_1/assets
```

In [14]:

```
pd.read_csv("Model1_results.csv", index_col = 'epoch')
```

Out[14]:

	accuracy	loss	val_accuracy	val_loss
epoch				
0	0.238084	2.478049	0.358722	2.173184
1	0.465777	1.879994	0.515880	1.739900
2	0.575512	1.595883	0.601602	1.531260
3	0.626775	1.462720	0.636637	1.437193
4	0.653899	1.378280	0.662754	1.354712
5	0.675194	1.313904	0.666758	1.329382
6	0.686002	1.268955	0.682683	1.270512
7	0.699525	1.229328	0.702976	1.213277
8	0.708887	1.196150	0.702248	1.202934
9	0.716018	1.168678	0.701793	1.186524
10	0.721173	1.146621	0.706161	1.175021
11	0.727244	1.126974	0.710802	1.157863
12	0.730841	1.109177	0.720994	1.137612
13	0.736253	1.091938	0.729366	1.115517
14	0.738710	1.079546	0.737101	1.080853
15	0.742083	1.063757	0.726454	1.111445
16	0.748940	1.049841	0.743107	1.061603
17	0.748908	1.041760	0.730276	1.092971
18	0.751172	1.030892	0.725726	1.108838
19	0.754416	1.021158	0.746019	1.040681
20	0.757066	1.013108	0.737192	1.076912
21	0.760937	1.002429	0.760397	1.011521
22	0.763715	0.994634	0.753117	1.020146
23	0.765819	0.987059	0.747839	1.035562
24	0.767312	0.978705	0.753481	1.021911
25	0.769721	0.972923	0.758031	1.003064
26	0.770942	0.965171	0.754209	1.019864
27	0.774587	0.957152	0.764310	0.991527
28	0.775390	0.952242	0.757576	1.002277
29	0.778345	0.947269	0.764037	0.977804


```

In [15]: fig = plt.figure(figsize=(12, 5))
fig.add_subplot(121)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Training Loss vs. epochs')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')

fig.add_subplot(122)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Training Accuracy vs. epochs')
plt.ylabel('Acc')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='lower right')

```

Out[15]: <matplotlib.legend.Legend at 0x7fabee302590>



```

In [16]: #Loss and accuracy on test set

test = model.evaluate(test_images_gr, test_labels, batch_size = 10, return_
print("Loss",test['loss'])
print("Accuracy", test['accuracy'])

```

2604/2604 [=====] - 2s 829us/step - loss: 1.0516 - accuracy: 0.7488  
Loss 1.051615834236145  
Accuracy 0.7487707734107971

## CNN neural network classifier



In [17]:

```
def CNN_model():
    model = Sequential([
        Conv2D(16,(3,3), padding='SAME', activation='relu', input_shape=(32,32,3)),
        BatchNormalization(),
        Dropout(.2),
        MaxPooling2D((3,3)),
        Conv2D(16,(3,3), padding='SAME', activation='relu'),
        BatchNormalization(),
        MaxPooling2D((3,3)),
        Flatten(),
        Dense(64, activation='relu',kernel_initializer='he_normal',
            kernel_regularizer=tf.keras.regularizers.l2(0.001)),
        Dense(10, activation='softmax')

    ])
    return model
```

In [18]:

```
print(train_images[:, :, :].shape)
model2 = CNN_model()
```

(73257, 32, 32)

In [19]:

```
model2.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 16)	160
batch_normalization (Batch Normalization)	(None, 32, 32, 16)	64
dropout (Dropout)	(None, 32, 32, 16)	0
max_pooling2d (MaxPooling2D)	(None, 10, 10, 16)	0
conv2d_1 (Conv2D)	(None, 10, 10, 16)	2320
batch_normalization_1 (Batch Normalization)	(None, 10, 10, 16)	64
max_pooling2d_1 (MaxPooling2D)	(None, 3, 3, 16)	0
flatten_1 (Flatten)	(None, 144)	0
dense_4 (Dense)	(None, 64)	9280
dense_5 (Dense)	(None, 10)	650
Total params: 12,538		
Trainable params: 12,474		
Non-trainable params: 64		

In [20]:

```
model2.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.0001),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy'])
print(train_images[0,:,:].shape)
print(train_labels.shape)
print(train_images.shape)

checkpoint = ModelCheckpoint('Model_2', save_weights_only=False, save_best_

history2 = model2.fit(train_images[...np.newaxis],train_labels, epochs=30,
                      callbacks=[checkpoint, tf.keras.callbacks.CSVLogger("Mc
```

```
(32, 32)
(73257, 1)
(73257, 32, 32)
Epoch 1/30
6227/6227 [=====] - 50s 8ms/step - loss: 1.9614 -
accuracy: 0.3953 - val_loss: 1.3333 - val_accuracy: 0.6132
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 2/30
6227/6227 [=====] - 49s 8ms/step - loss: 1.0951 -
accuracy: 0.6896 - val_loss: 1.0029 - val_accuracy: 0.7195
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 3/30
6227/6227 [=====] - 50s 8ms/step - loss: 0.8861 -
accuracy: 0.7607 - val_loss: 0.8737 - val_accuracy: 0.7675
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 4/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.7951 -
accuracy: 0.7877 - val_loss: 0.8059 - val_accuracy: 0.7904
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 5/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.7374 -
accuracy: 0.8050 - val_loss: 0.7717 - val_accuracy: 0.7980
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 6/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.6962 -
accuracy: 0.8153 - val_loss: 0.7275 - val_accuracy: 0.8055
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 7/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.6633 -
accuracy: 0.8261 - val_loss: 0.6842 - val_accuracy: 0.8242
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 8/30
6227/6227 [=====] - 50s 8ms/step - loss: 0.6357 -
accuracy: 0.8336 - val_loss: 0.6771 - val_accuracy: 0.8206
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 9/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.6124 -
accuracy: 0.8398 - val_loss: 0.6440 - val_accuracy: 0.8298
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 10/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.5937 -
accuracy: 0.8426 - val_loss: 0.6336 - val_accuracy: 0.8367
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 11/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.5744 -
accuracy: 0.8490 - val_loss: 0.6120 - val_accuracy: 0.8378
```

```
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 12/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.5574 -
accuracy: 0.8535 - val_loss: 0.6215 - val_accuracy: 0.8340
Epoch 13/30
6227/6227 [=====] - 48s 8ms/step - loss: 0.5407 -
accuracy: 0.8578 - val_loss: 0.6209 - val_accuracy: 0.8317
Epoch 14/30
6227/6227 [=====] - 50s 8ms/step - loss: 0.5305 -
accuracy: 0.8600 - val_loss: 0.5871 - val_accuracy: 0.8440
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 15/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.5185 -
accuracy: 0.8633 - val_loss: 0.5701 - val_accuracy: 0.8489
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 16/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.5060 -
accuracy: 0.8655 - val_loss: 0.5630 - val_accuracy: 0.8481
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 17/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.4932 -
accuracy: 0.8693 - val_loss: 0.5613 - val_accuracy: 0.8498
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 18/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.4851 -
accuracy: 0.8717 - val_loss: 0.5370 - val_accuracy: 0.8590
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 19/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.4741 -
accuracy: 0.8743 - val_loss: 0.5447 - val_accuracy: 0.8537
Epoch 20/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.4651 -
accuracy: 0.8752 - val_loss: 0.5360 - val_accuracy: 0.8559
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 21/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.4569 -
accuracy: 0.8789 - val_loss: 0.5324 - val_accuracy: 0.8564
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 22/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.4483 -
accuracy: 0.8794 - val_loss: 0.5241 - val_accuracy: 0.8599
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 23/30
6227/6227 [=====] - 48s 8ms/step - loss: 0.4413 -
accuracy: 0.8828 - val_loss: 0.5130 - val_accuracy: 0.8601
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 24/30
6227/6227 [=====] - 48s 8ms/step - loss: 0.4350 -
accuracy: 0.8832 - val_loss: 0.5110 - val_accuracy: 0.8614
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 25/30
6227/6227 [=====] - 48s 8ms/step - loss: 0.4317 -
accuracy: 0.8830 - val_loss: 0.5109 - val_accuracy: 0.8600
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 26/30
6227/6227 [=====] - 48s 8ms/step - loss: 0.4241 -
accuracy: 0.8857 - val_loss: 0.4922 - val_accuracy: 0.8671
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 27/30
6227/6227 [=====] - 48s 8ms/step - loss: 0.4213 -
accuracy: 0.8865 - val_loss: 0.4914 - val_accuracy: 0.8698
```

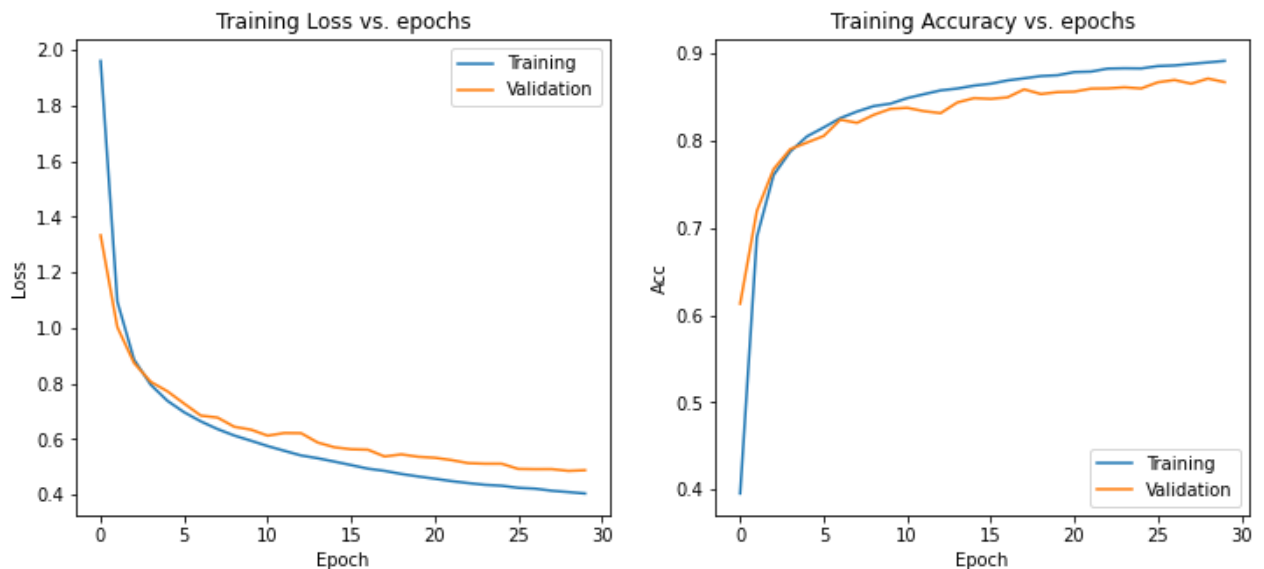
```
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 28/30
6227/6227 [=====] - 48s 8ms/step - loss: 0.4136 -
accuracy: 0.8883 - val_loss: 0.4915 - val_accuracy: 0.8655
Epoch 29/30
6227/6227 [=====] - 48s 8ms/step - loss: 0.4088 -
accuracy: 0.8900 - val_loss: 0.4851 - val_accuracy: 0.8714
INFO:tensorflow:Assets written to: Model_2/assets
Epoch 30/30
6227/6227 [=====] - 49s 8ms/step - loss: 0.4037 -
accuracy: 0.8917 - val_loss: 0.4875 - val_accuracy: 0.8673
```

In [21]:

```
fig = plt.figure(figsize=(12, 5))
fig.add_subplot(121)
plt.plot(history2.history['loss'])
plt.plot(history2.history['val_loss'])
plt.title('Training Loss vs. epochs')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')

fig.add_subplot(122)
plt.plot(history2.history['accuracy'])
plt.plot(history2.history['val_accuracy'])
plt.title('Training Accuracy vs. epochs')
plt.ylabel('Acc')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='lower right')
```

Out[21]: <matplotlib.legend.Legend at 0x7fabed75f610>



In [22]:

```
#Loss and accuracy on test set

test = model2.evaluate(test_images_gr[...], np.newaxis], test_labels, batch_size)
print("Loss", test['loss'])
print("Accuracy", test['accuracy'])
```

```
2604/2604 [=====] - 6s 2ms/step - loss: 0.5318 - accuracy: 0.8535
Loss 0.5318433046340942
Accuracy 0.8535264134407043
```

# Models predictions

In [23]:

```
#Loading best weights and models for each version

MLP = load_model('Model_1')
CNN = load_model('Model_2')
```

In [24]:

```
#Randomly select 5 images and labels from the test set

random_ix = np.random.choice(test_images.shape[3], 5) #random index to print
random_im = test_images[:, :, :, random_ix] #take images to print
random_labels = test_labels[random_ix] #take labels
random_images_gr = convert_gray(random_im, random_labels, False)
```




In [45]:

```
#Results




predictionsMLP = MLP.predict(random_images_gr)
predictionsCNN = CNN.predict(random_images_gr[:, :, :, np.newaxis])

fig, axes = plt.subplots(5,3, figsize=(14, 10))
fig.subplots_adjust(hspace=.8, wspace=0.2)
for (i, ax) in enumerate(axes):
    axes[i, 0].set_axis_off()
    axes[i, 0].imshow(np.squeeze(random_im[:, :, :, i]))
    axes[i, 0].set_title('Test Label: ' +str(random_labels[i]), fontsize=10)
    axes[i, 1].bar(np.arange(len(predictionsMLP[i])), predictionsMLP[i])
    axes[i, 1].set_title(f"MLP Model's predictive \n distribution. Predictions")
    axes[i, 2].bar(np.arange(len(predictionsCNN[i])), predictionsCNN[i])
    axes[i, 2].set_title(f"CNN Model's predictive \n distribution. Predictions")
```




Test Label: [9]






Test Label: [5]



Test Label: [5]



Test Label: [2]



Test Label: [9]

